

# Performance Tuning in Production

DEVVOX<sup>TM</sup>  
United Kingdom

James Gough  
Sadiq Jaffer  
Kirk Pepperdine  
Richard Warburton



# Session Overview

- Optimizing Java: a brief tour of the JVM
- Moving to G1GC
- Production Profiling: What, Why and How



# Optimizing Java: A JVM Tour

DEVVOXX<sup>TM</sup>  
United Kingdom

James Gough

@Jim\_\_Gough

<http://jamesgough.net>





# This Talk


- Who Am I
- Creating Bytecode
- Classloading
- Profiling Code
- Runtime Optimisations
- JITWatch



# About Me



- Started programming BASIC on the C64
- Worked as a Java and Web Developer
- Helped to design and test JSR 310
- Spent 4 years training Java and C++
- Written a book called Optimizing Java
- Work at Morgan Stanley
  - Building Client Facing Technology






# Creating Bytecode



+



=

JVM



# Creating Bytecode



+



=


**Java  
source  
code**

JVM





# Creating Bytecode




Class file creation





# Creating Bytecode



Class file creation



# The anatomy of a classfile

<b>M</b> agic Number	0xCAFEBAFE
<b>V</b> ersion of Class File Format	The minor and major versions of the class file
<b>C</b> onstant Pool	Pool of constants for the class
<b>A</b> ccess Flags	For example whether the class is abstract, static, etc.
<b>T</b> his Class	The name of the current class
<b>S</b> uper Class	The name of the super class
<b>I</b> nterfaces	Any interfaces in the class
<b>F</b> ields	Any fields in the class
<b>M</b> ethods	Any methods in the class
<b>A</b> tttributes	Any attributes of the class (e.g. name of the sourcefile, etc.)



# The anatomy of a classfile



My	Very	Cute	Animal	Turns	Savage	In	Full	Moon	Areas
M	V	C	A	T	S	I	F	M	A
Magic	Version	Constant	Access	This	Super	Interfaces	Fields	Methods	Attributes

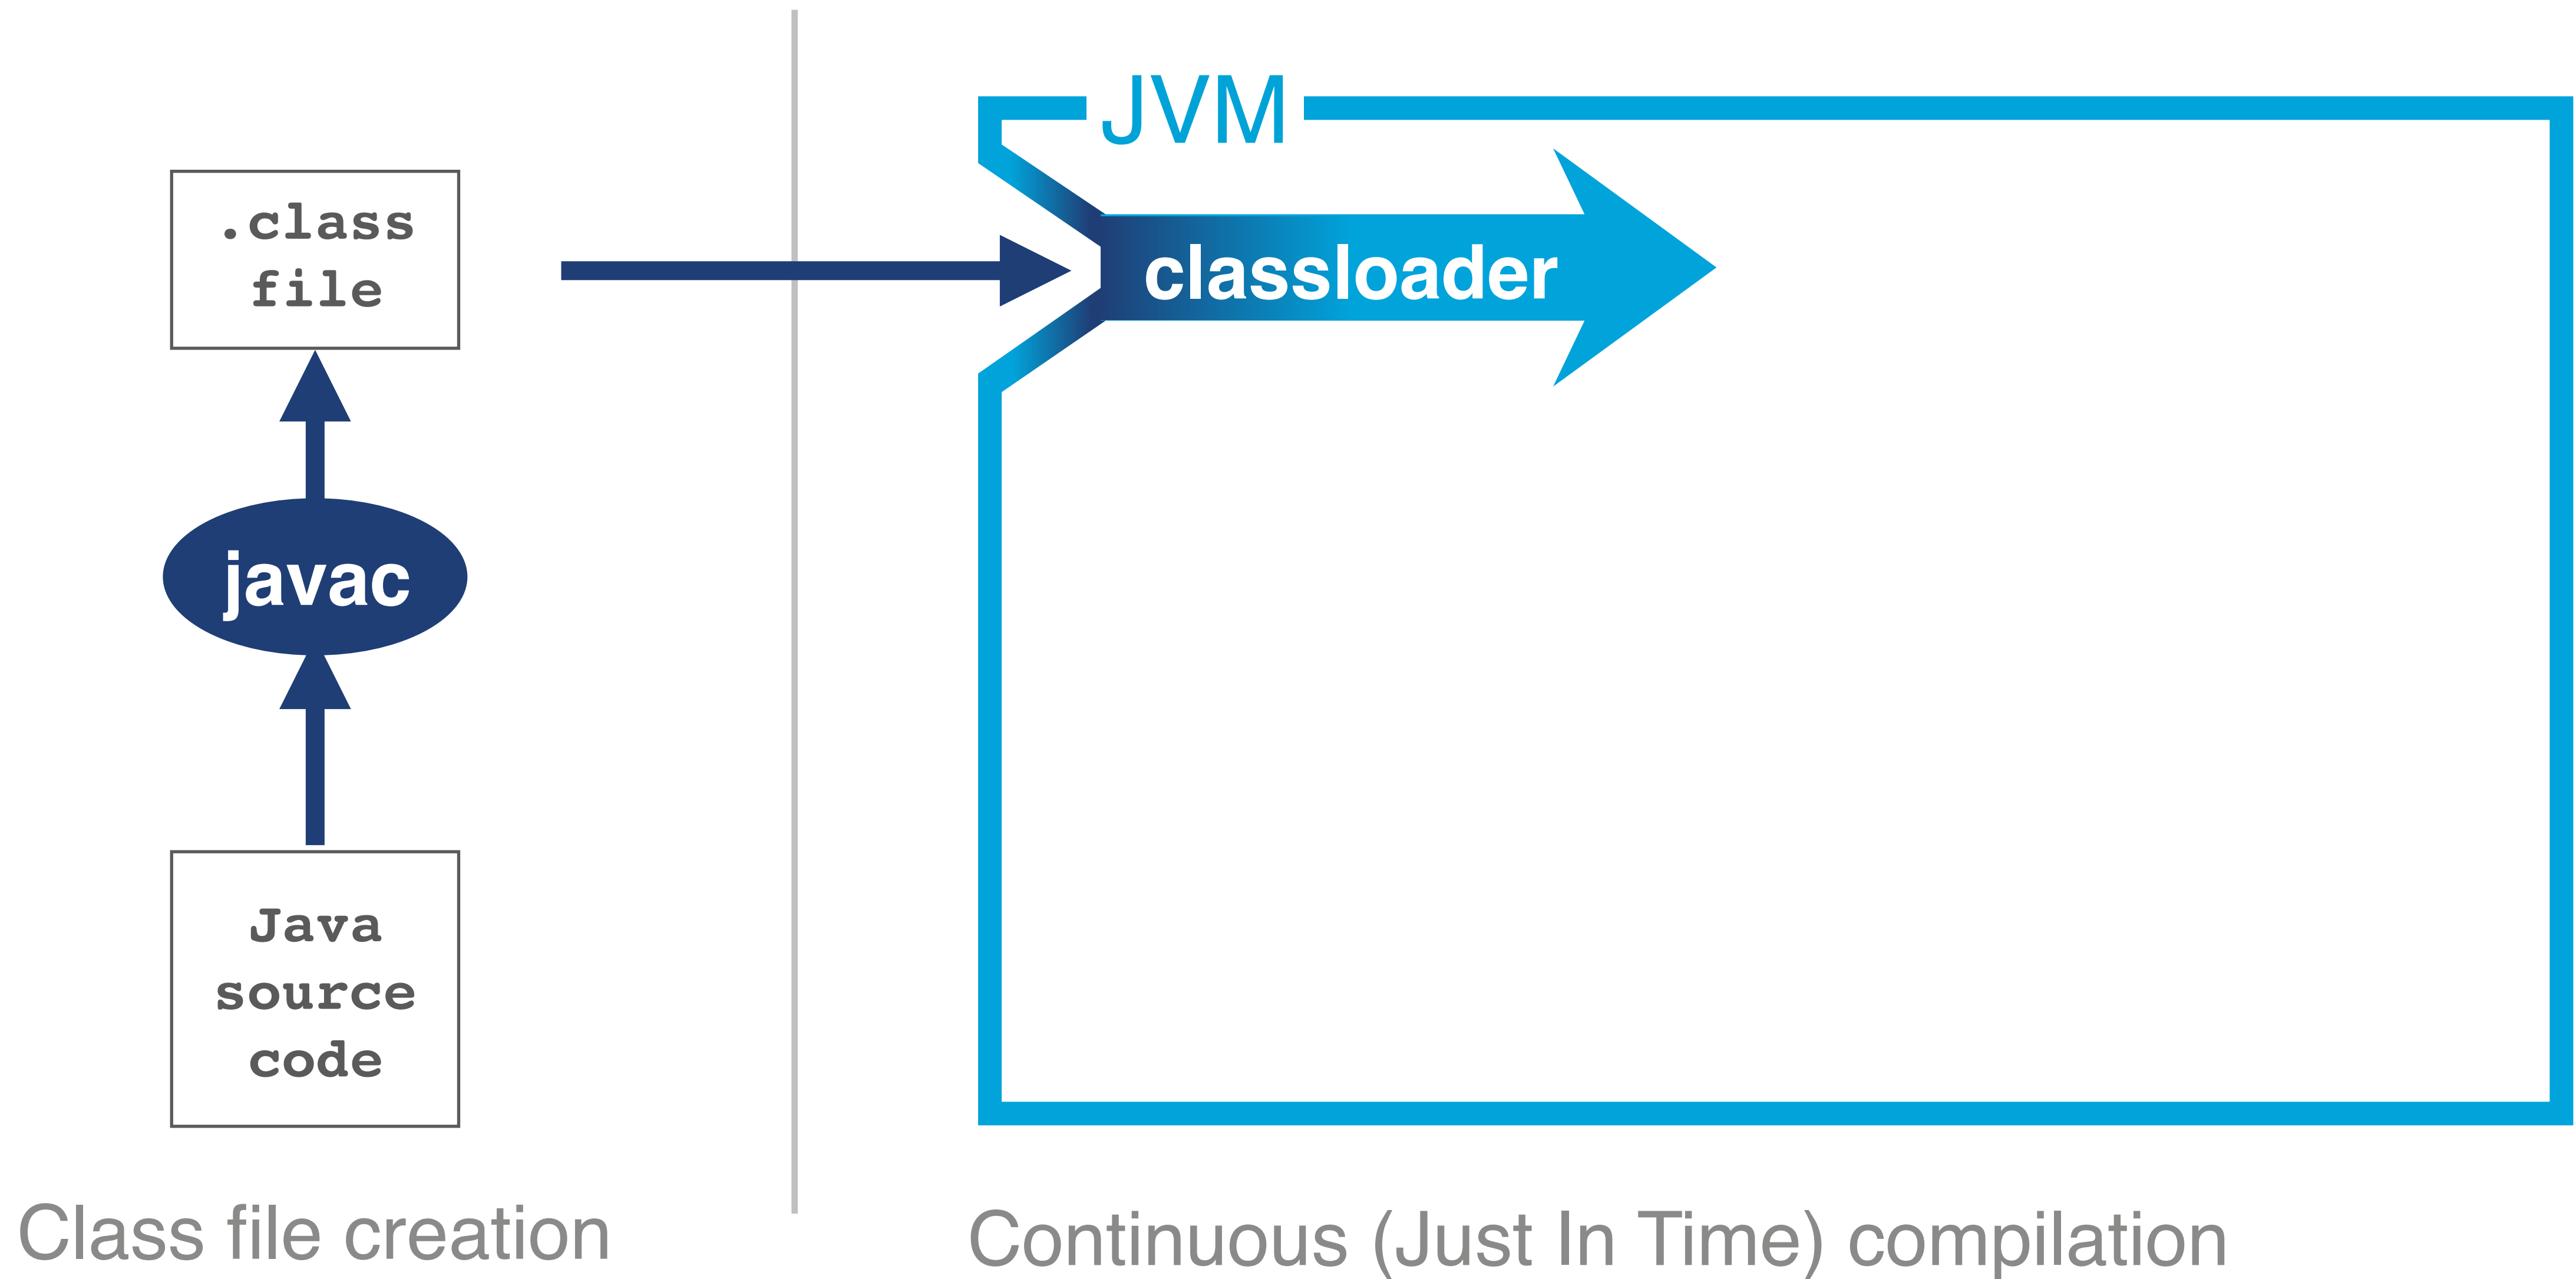
# Type Descriptors

- Describe signatures
- Common in `javap` output
- E.g.
  - `()Ljava/lang/String;`
  - `(I)V`
  - `(Ljava/lang/String;I)J`

Descriptor	Type
B	byte
C	char
D	double
F	float
I	int
J	long
L<type>;	Reference type
S	short
Z	boolean
[	Array-of



# How Bytecode is executed

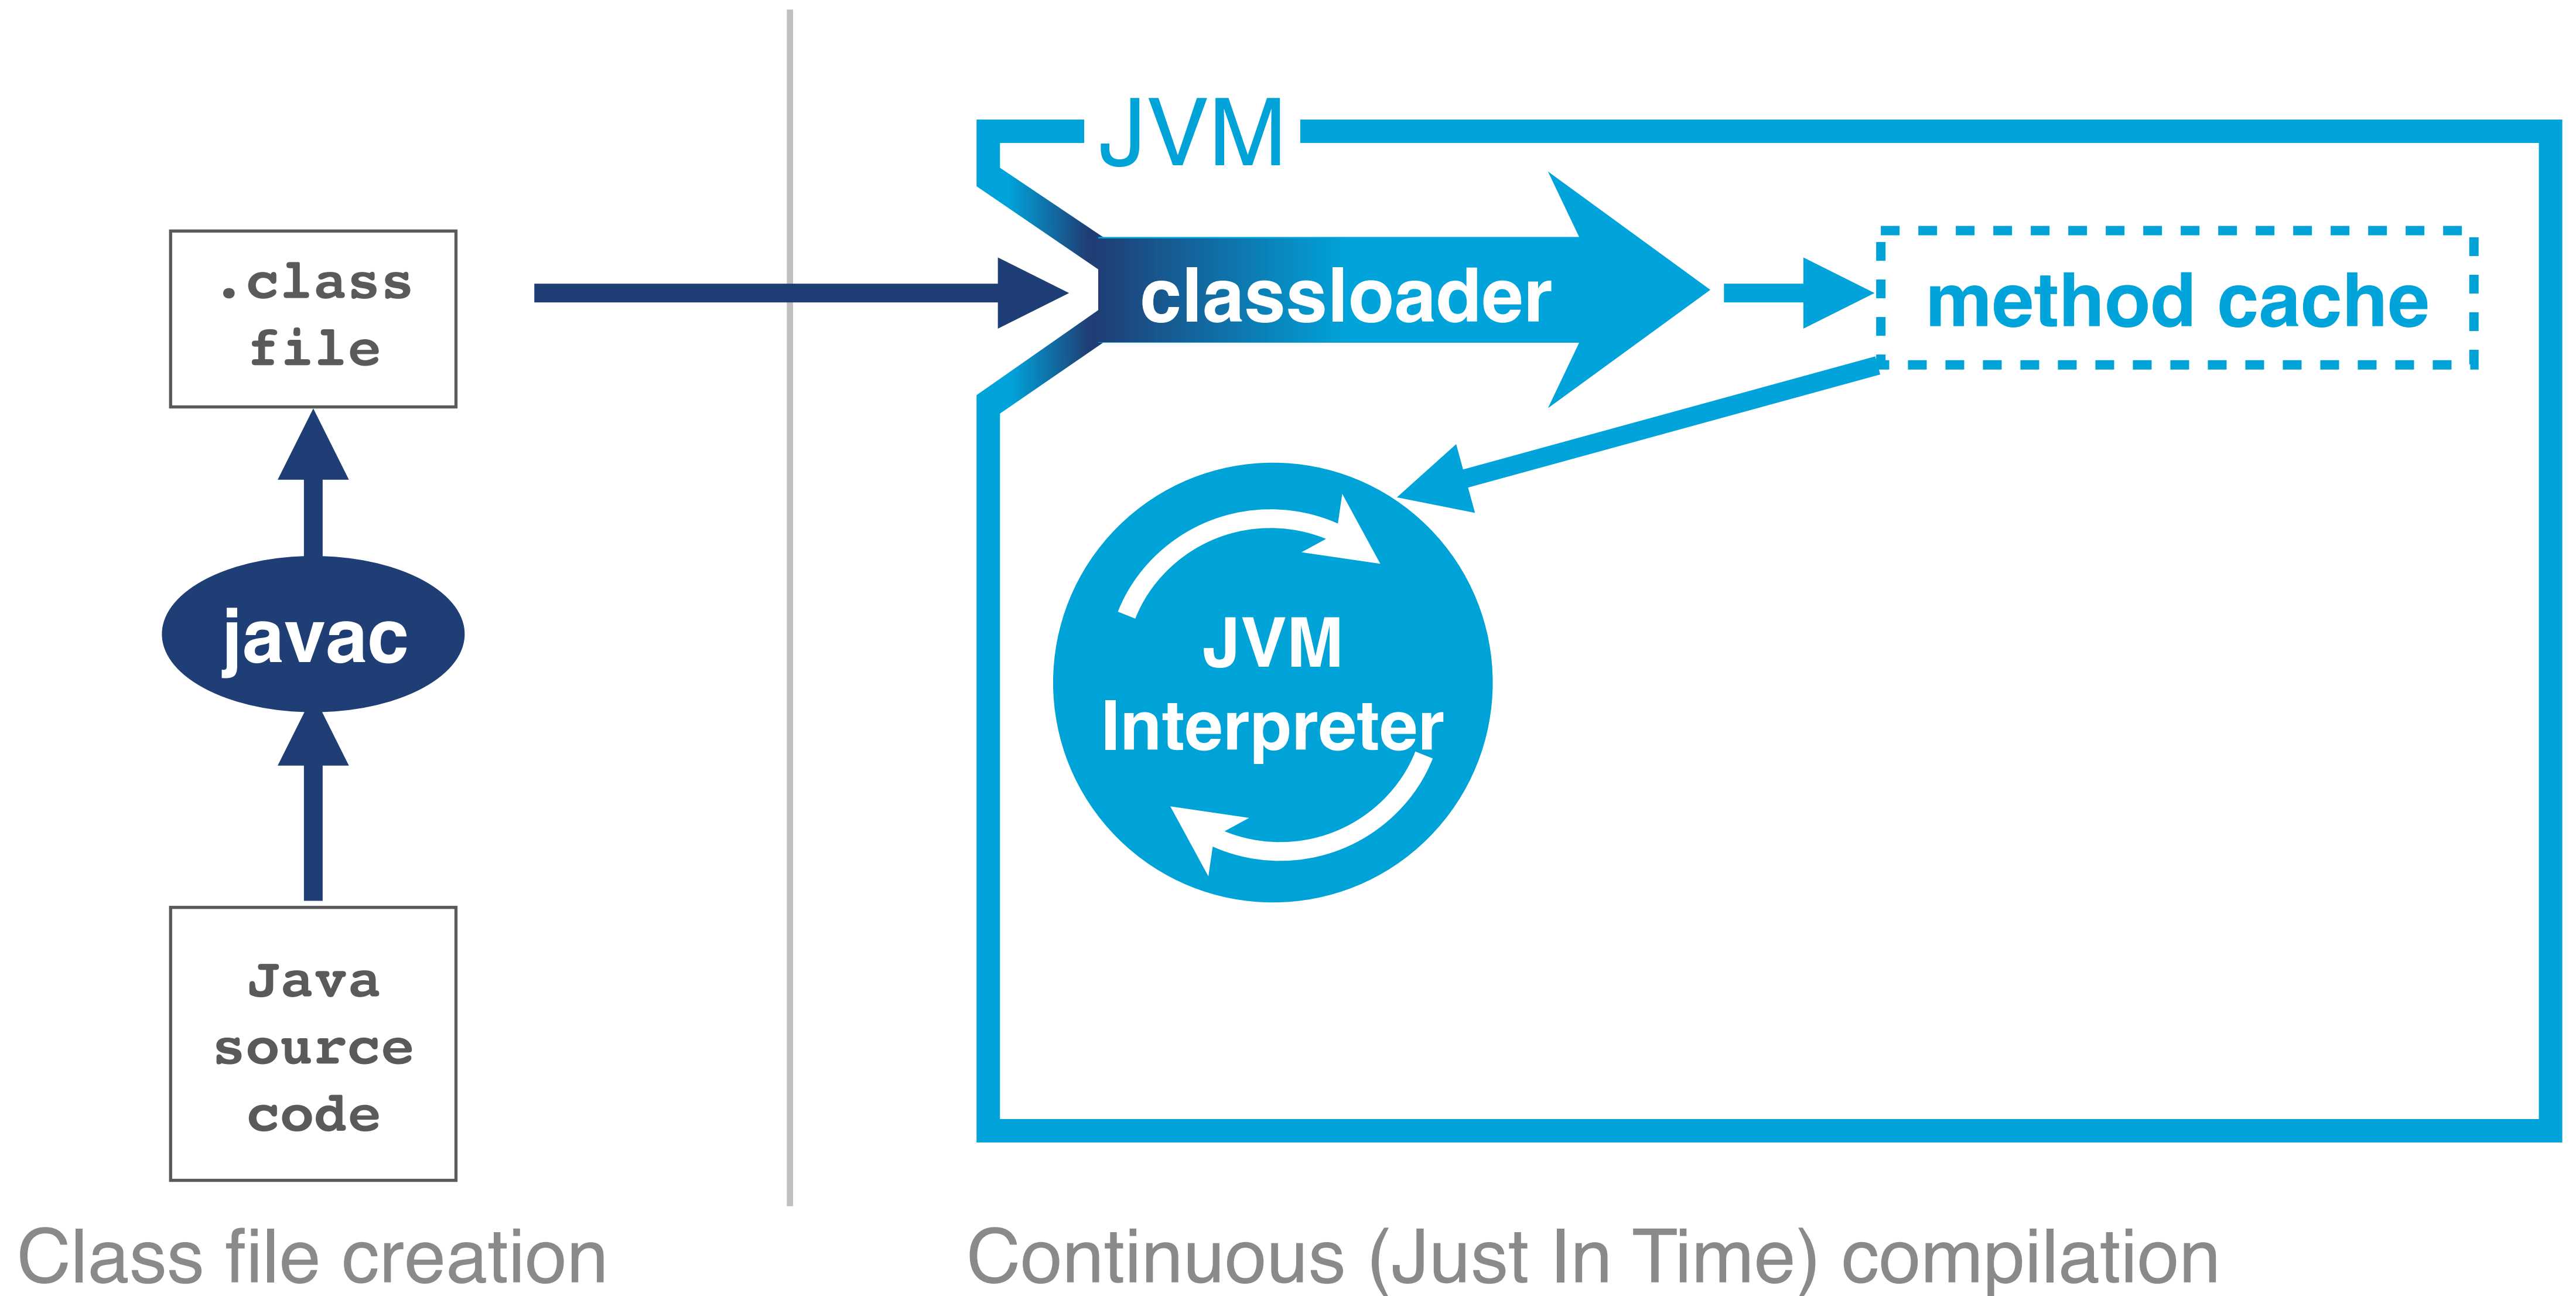


# Classloaders

- Classes are loaded just before they are needed
  - proven by the painful `ClassNotFoundException`
- Loads classfile into the `Class` object
  - mechanism for representing classes in the VM
- Example used in **watching-classloader**
  - <https://github.com/jpgough/watching-classloader>



# Interpreting Bytecode



# Interpreting Bytecode

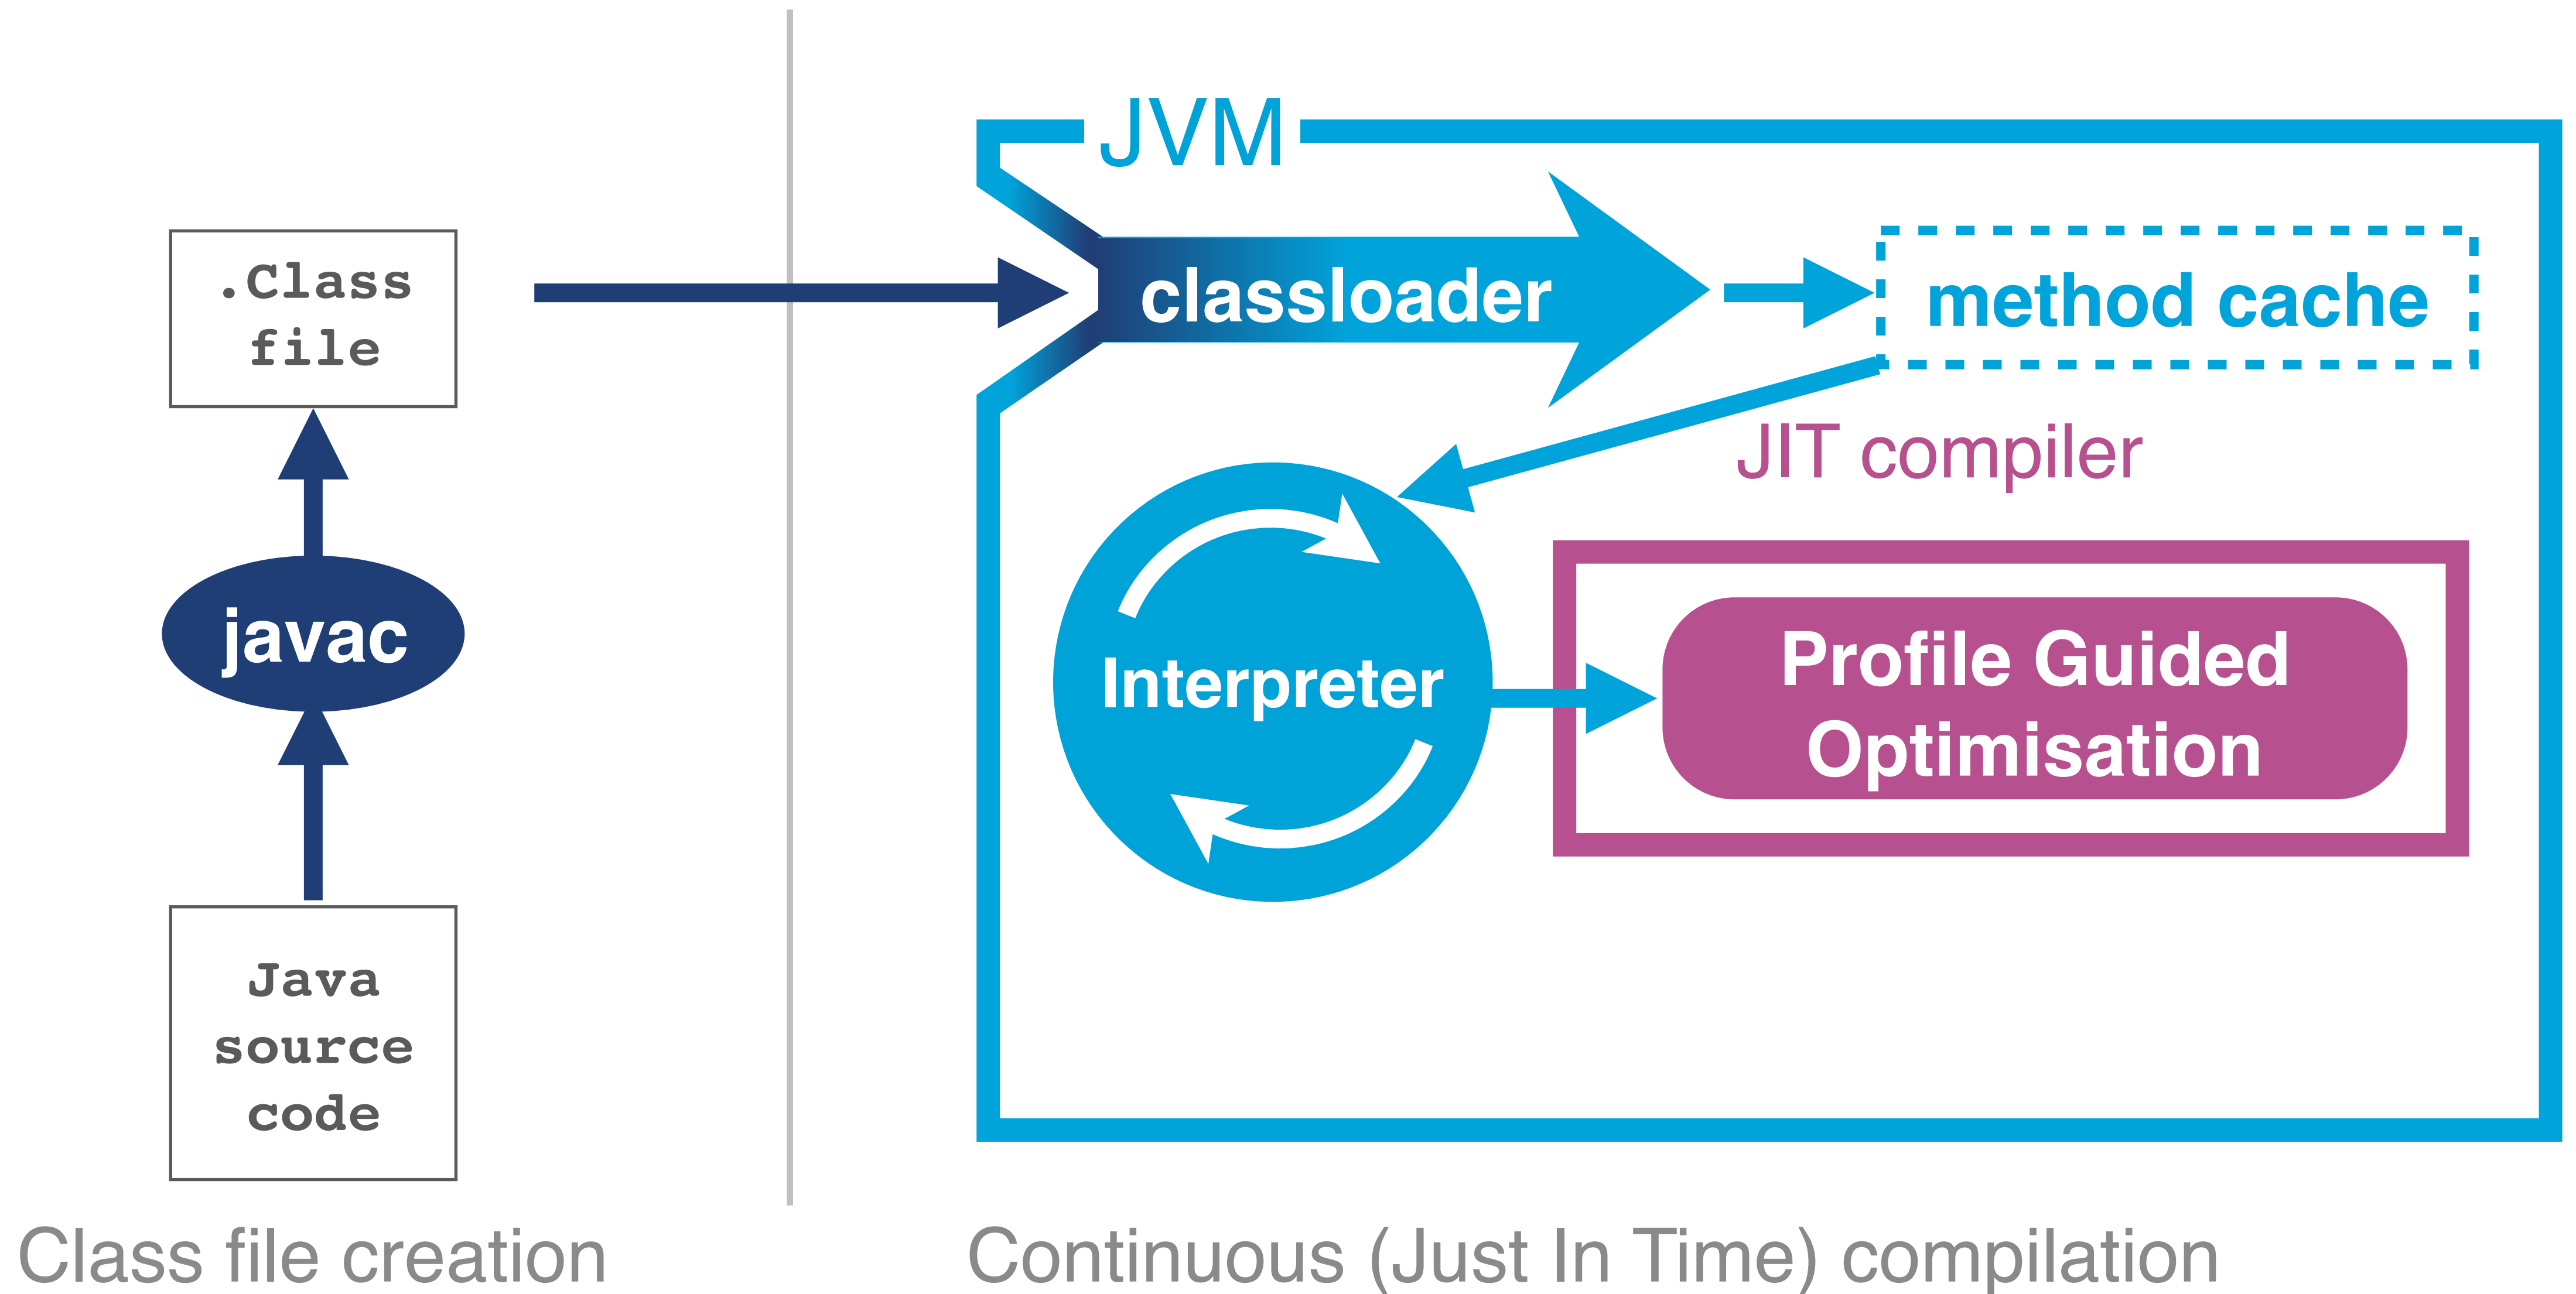
- Bytecode initially fully interpreted
- Conversion of each instruction to machine instruction
- Time not spent compiling code that is only used once



# How does Interpreting Help?

- Provides the opportunity to observe code execution paths
  - may not be the same for each execution of the app
- The profiler observes the execution and looks for the best optimisations
- Code is compiled after hitting a threshold
  - Configurable
  - JVM can revert optimisation decisions

# Profiling Code

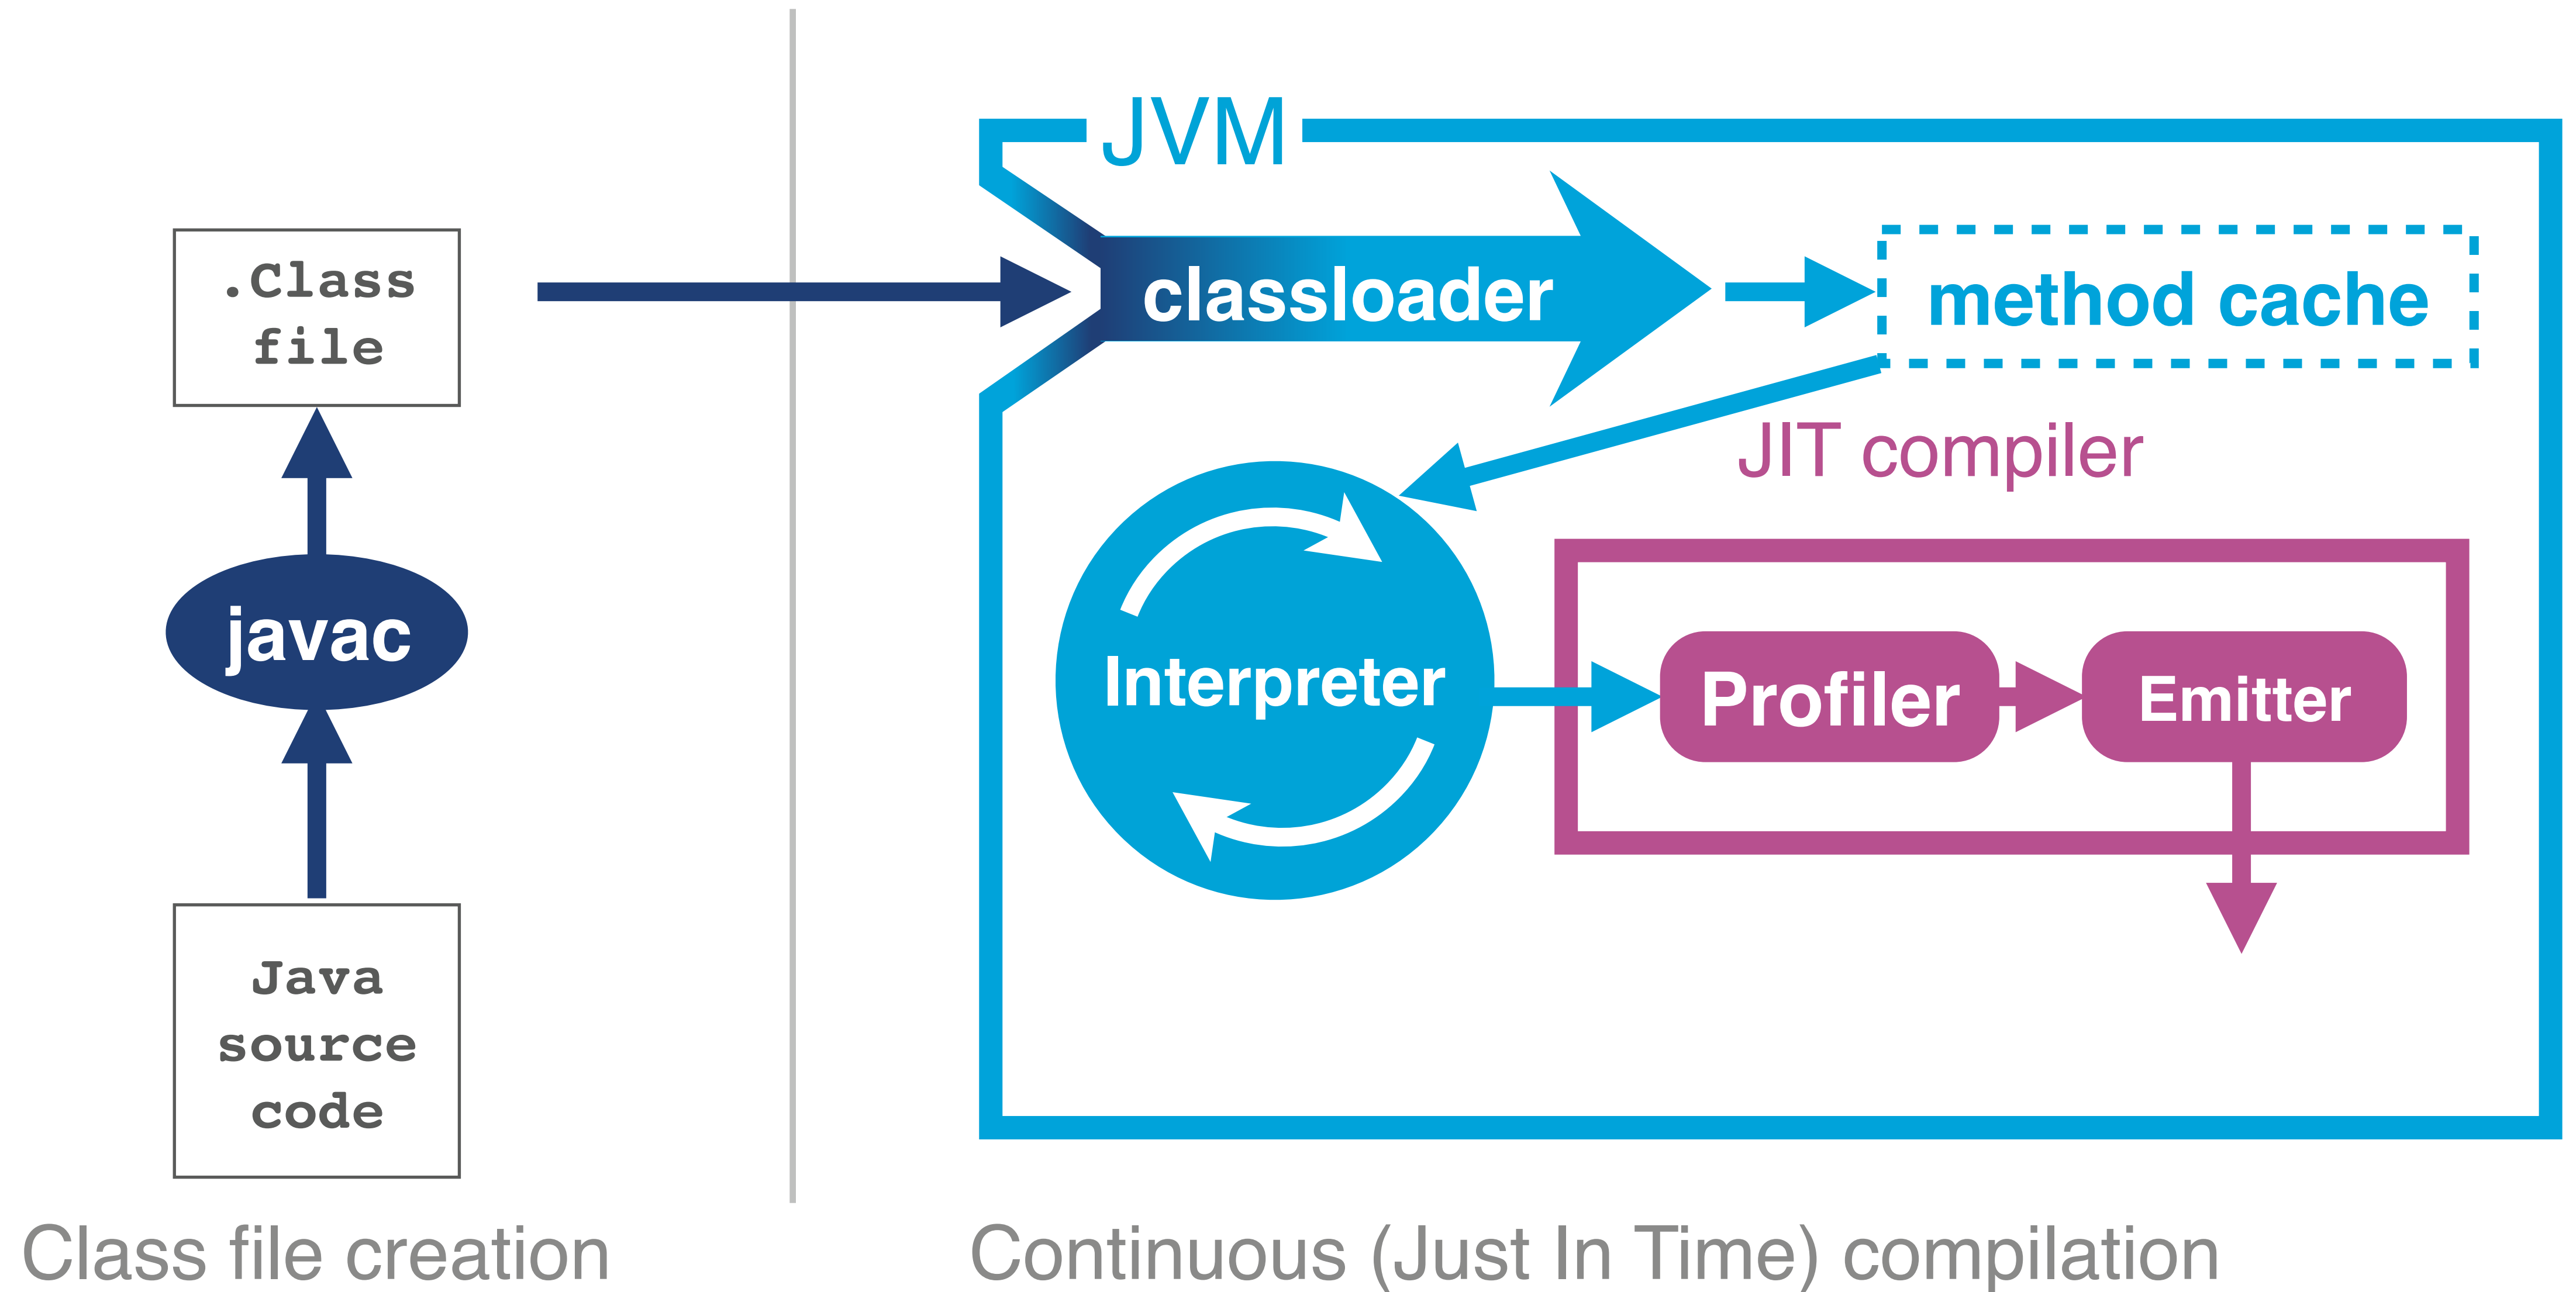




# Profiling Code

- Looking for loops or frequent execution of code blocks
- Barometer used to count the number of executions
- Threshold is reached and mode changes to tracing
- Tracing follows the execution path involving that method
  - proactively looking for optimisation opportunities
  - often stored as an intermediate representation
  - traces are used in the code generation phase

# The Hotspot JVM





# Viewing Code Compilation

```
java -XX:-TieredCompilation -XX:+PrintCompilation HelloWorld 2> /dev/null
```

Time Offset	Task	Method Name (size of compiled code)
321	40	sun.nio.cs.StreamEncoder::isOpen (5 bytes)
322	41	sun.nio.cs.StreamEncoder::implFlushBuffer (15 bytes)
327	42	sun.nio.cs.StreamEncoder::writeBytes (132 bytes)
331	43 !	java.io.PrintStream::write (69 bytes)
335	44 s	java.io.BufferedOutputStream::write (67 bytes)
337	46	java.nio.Buffer::clear (20 bytes)
337	47	java.lang.String::indexOf (7 bytes)
338	48 !	java.io.PrintStream::println (24 bytes)
338	49	java.io.PrintStream::print (13 bytes)
343	50 !	java.io.PrintStream::write (83 bytes)
346	51 !	java.io.PrintStream::newLine (73 bytes)
347	52	java.io.BufferedWriter::newLine (9 bytes)
347	53 %	HelloWorld::main @ 2 (23 bytes)

! method has exception handler(s)  
s method declared synchronized  
n native method (no compilation, generate wrapper)  
% on-stack replacement used

# Inlining

- Calling a method has an overhead
  - creation of a new stack frame
  - copying values required to the stack frame
  - returning from the stack frame post execution
- Consider a method call in a for loop

```
public class HelloWorld {  
    public static void main(String[] args) {  
        for(int i=0; i < 100_000; i++) {  
            System.err.println("Hello World");  
        }  
    }  
}
```



# Inlining

```
java -XX:+UnlockDiagnosticVMOptions -XX:+PrintInlining HelloWorld 2> /dev/null
```

```
@ 40  java.io.BufferedOutputStream::flush (12 bytes)  inline (hot)
      \-> TypeProfile (19272/19272 counts) = java/io/BufferedOutputStream
      @ 1  java.io.BufferedOutputStream::flushBuffer (29 bytes)  inline (hot)
      @ 20  java.io.FileOutputStream::write (12 bytes)  inline (hot)
      \-> TypeProfile (4696/4696 counts) = java/io/FileOutputStream
      @ 8  java.io.FileOutputStream::writeBytes (0 bytes)  native method
      @ 8  java.io.OutputStream::flush (1 bytes)  inline (hot)
      \-> TypeProfile (7047/7047 counts) = java/io/FileOutputStream
!m    @ 13  java.io.PrintStream::println (24 bytes)
      @ 6  java.io.PrintStream::print (13 bytes)
!m    @ 9  java.io.PrintStream::write (83 bytes)  callee is too large
!m    @ 10  java.io.PrintStream::newLine (73 bytes)  callee is too large
!m    @ 13  java.io.PrintStream::println (24 bytes)
      @ 6  java.io.PrintStream::print (13 bytes)
!m    @ 9  java.io.PrintStream::write (83 bytes)  callee is too large
!m    @ 10  java.io.PrintStream::newLine (73 bytes)  callee is too large
!m    @ 13  java.io.PrintStream::println (24 bytes)  already compiled into a big method
```

# Constant Subexpression Elimination

- Compiler hunts through code for common expressions
  - if results analyses replacement with a single variable
- Relies on data flow analysis of the program
  - which is done during the profiling and tracing part



# Dead Code Elimination

- Removes code that is never executed
  - shrinks the size of the program
  - avoid executing irrelevant operations
- Dynamic dead code elimination
  - eliminated based on possible set of values
  - determined at runtime

# Register Allocation

- Identification of variables suitable for registers
  - to avoid cache misses
  - improve execution speed of the program
- Uses data from the trace to make informed decision



# Loop-Invariant Code Motion

- Involves removal of code from loops
  - for code that doesn't impact the outcome of the loop
  - moved above the loop to avoid unnecessary execution
- Hoisted code can now be cached in a register
  - improving performance of the loop execution


# Escape Analysis

- Introduced in later versions of Java 6
- Analyses code to assert if an object reference
  - returns or leaves the scope of the method
  - stored in global variables
- Allocates unescaped objects on the stack
  - avoids the cost of garbage collection
  - prevents workload pressures on Eden
  - beneficial effects to counter high infant mortality GC impact



# Loop Unrolling

```
private static final String[] RESPONSES =  
    { "Yes", "No", "Maybe" };  
  
public void processResponses () {  
    for ( String response: RESPONSES ) {  
        process(response);  
    }  
}
```



```
private static final String[] RESPONSES =  
    { "Yes", "No", "Maybe" };  
  
public void processResponses () {  
    process(RESPONSES[0]);  
    process(RESPONSES[1]);  
    process(RESPONSES[2]);  
}
```

# Loop Unrolling

@Benchmark

```
public long intStride1()
{
    long sum = 0;
    for (int i = 0; i < MAX; i++)
    {
        sum += data[i];
    }
    return sum;
}
```

@Benchmark

```
public long longStride1()
{
    long sum = 0;
    for (long l = 0; l < MAX; l++)
    {
        sum += data[(int) l];
    }
    return sum;
}
```

Benchmark	Mode	Cnt	Score	Error	Units
LoopUnrollingCounter.intStride1	thrpt	200	<b>2423.818</b>	± 2.547	ops/s
LoopUnrollingCounter.longStride1	thrpt	200	<b>1469.833</b>	± 0.721	ops/s

Excerpt From: Benjamin J. Evans, James Gough, and Chris Newland. “Optimizing Java.” iBooks.



# Loop Unrolling

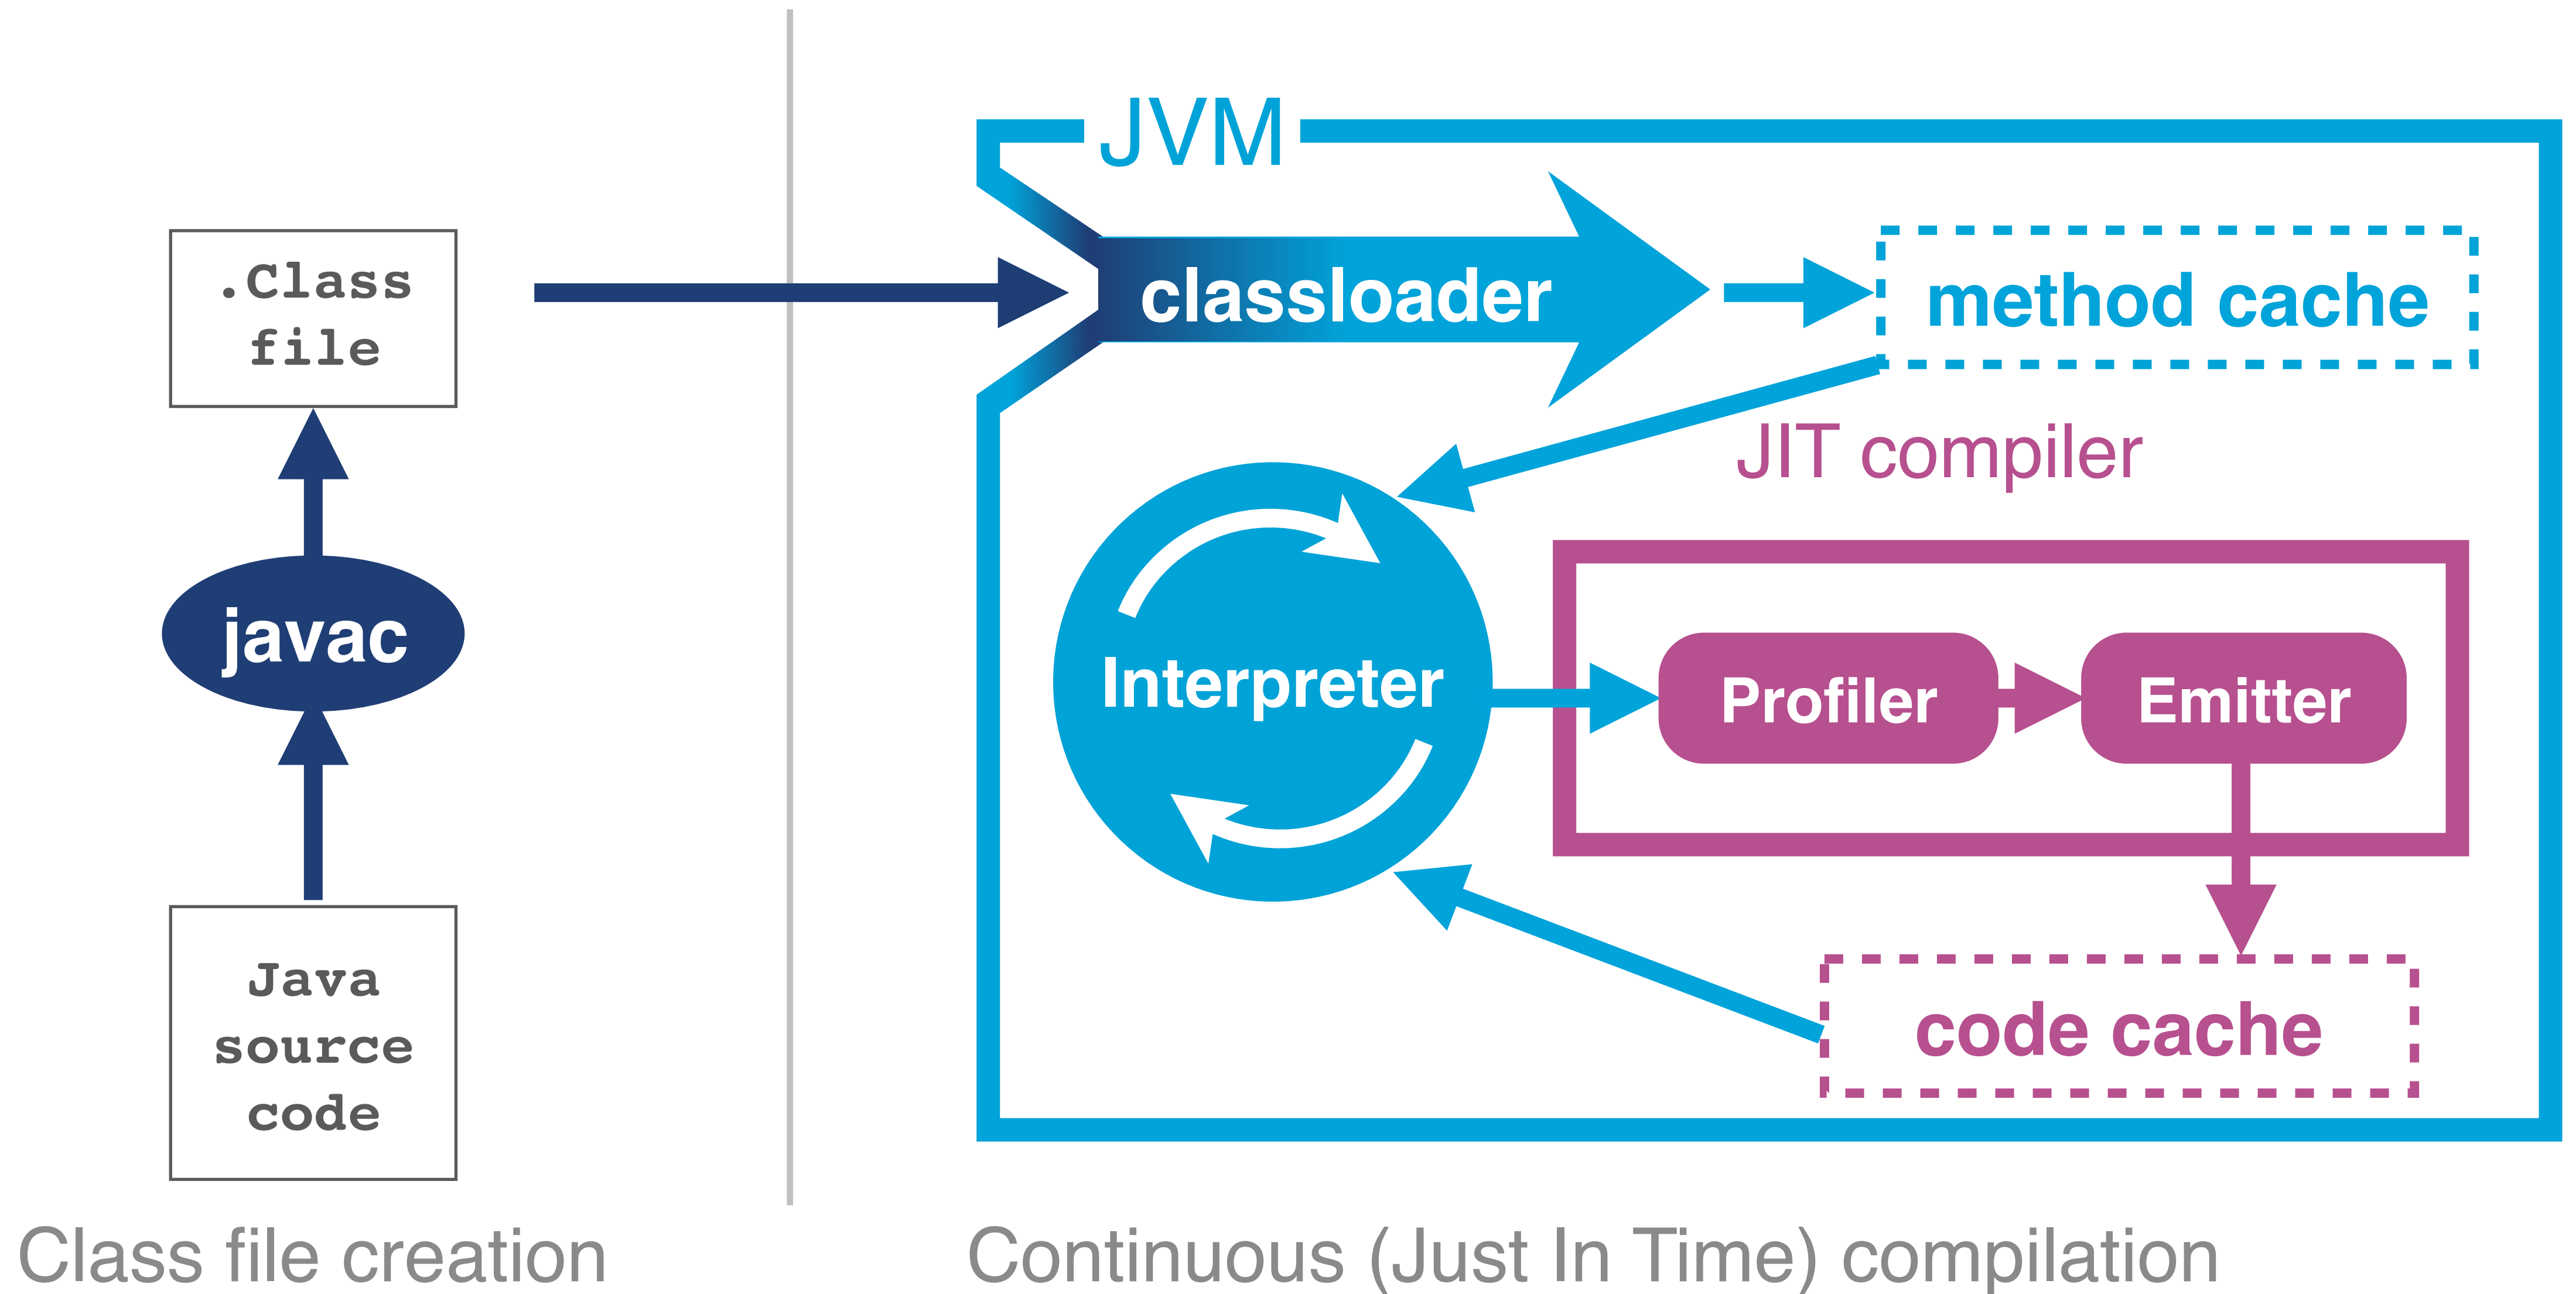
- Can unroll int, char and short loops
- Can remove safe point checks
- Removes back branches and branch prediction cost
- Reduces the work needed by each “iteration”

# Monomorphic Dispatch

- When HotSpot encounters a virtual call site, often only one type will ever be seen there
  - e.g. There's only one implementing class for an interface
- Hotspot can optimize vtable lookup
  - Subclasses have the same vtable structure as their parent
  - Hotspot can collapse the child into the parent
- Classloading tricks can invalidate monomorphic dispatch
  - The class word in the header is checked
  - If changed then this optimisation is backed out



# Code Cache

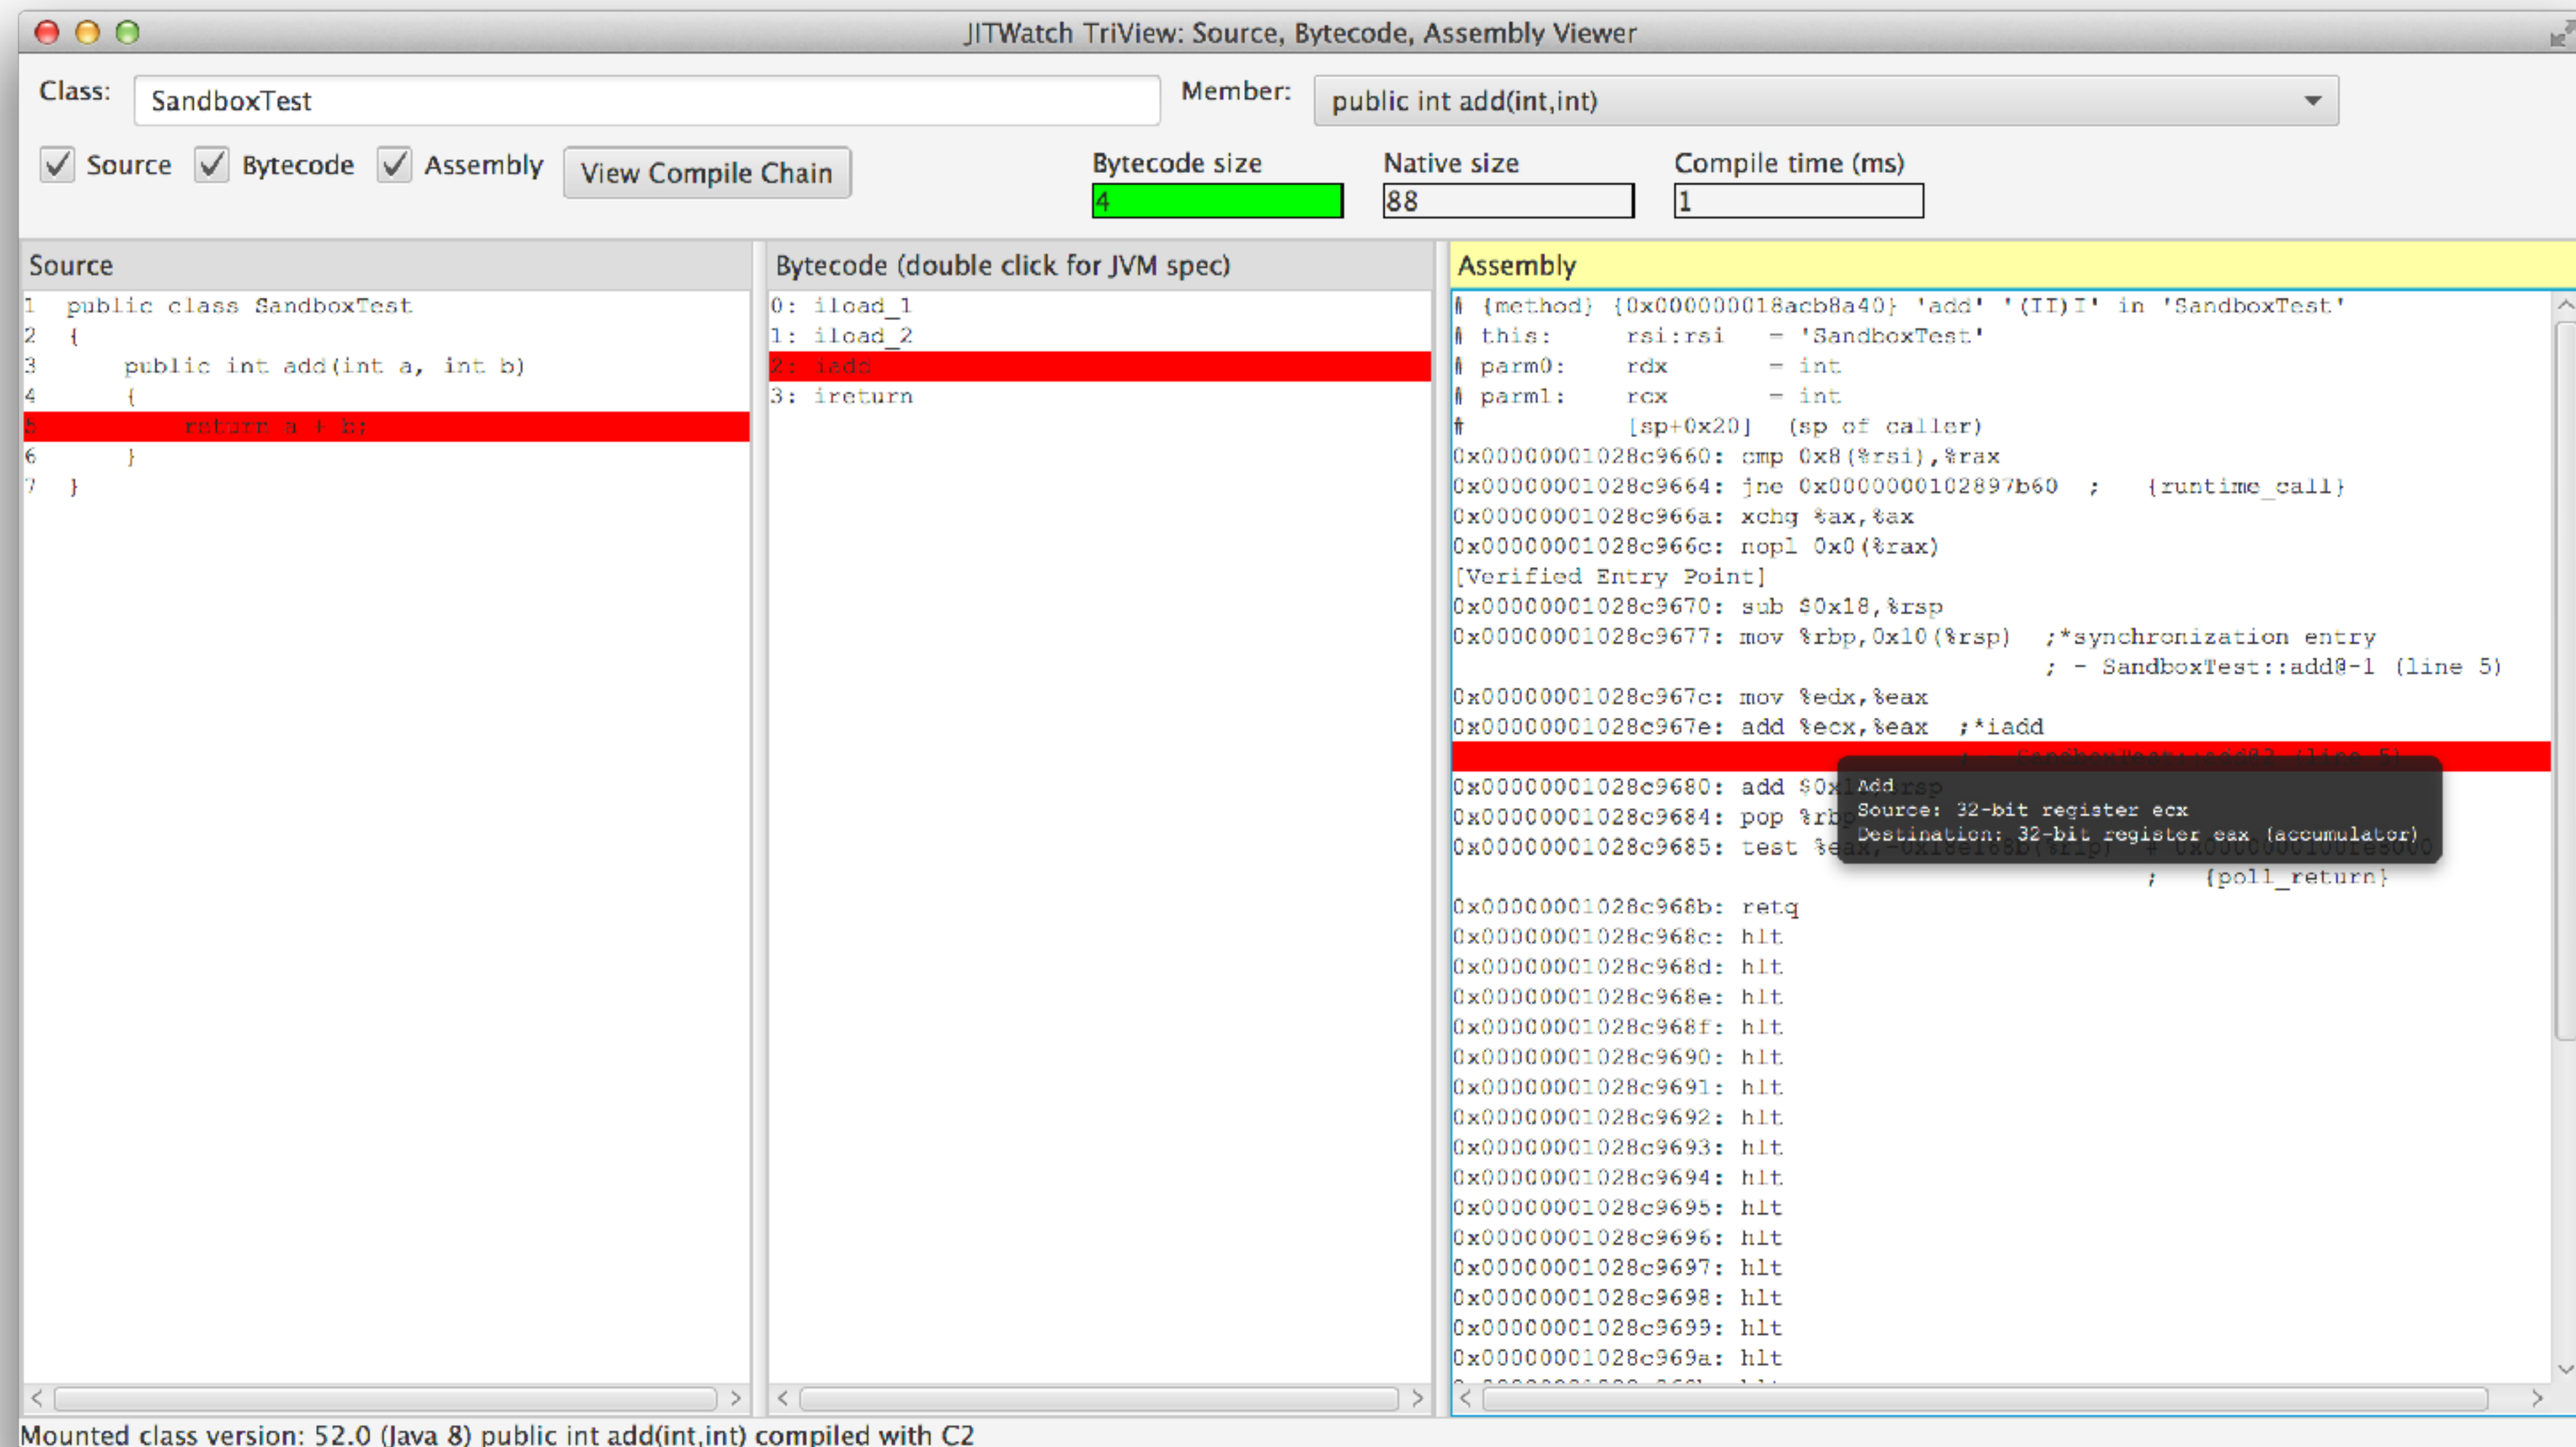


# Code Cache

- The code cache contains the JIT native compiled code
- Code is JIT'd on a per method basis
  - 1. This occurs when an entry counter is exceeded
  - 2. Internal Representation (IR) is built
  - 3. Optimisations are applied
  - 4. JIT turns IR into native code
- Pointers are swizzled to use the native code
  - native code is executed on the next call



# Introduction to JITWatch



# Summary

- Java has carried a brand name of being slow
- Java can emit instructions comparable to C++
- javac doesn't do much optimisation
- We can make better decisions from profiling at runtime
- JITWatch makes life easier



# Performance Landscape

